

IMPLEMENTING A PIPELINED CPU

Unit 2


1

PIPELINED Y86 IMPLEMENTATION


- Unit outline
 - Motivation and basic concepts
 - Initial Implementation
 - Hazards
 - Types of hazards
 - Dealing with hazards by stalling
 - Data hazards: forwarding to avoid stalling
 - Control hazards: branch prediction
 - Indirect jumps
 - Performance analysis

MOTIVATION

In a sequential Y86 implementation



HUMAN PIPELINE



PIPELINE COMPUTATION

- Represent each stage as a module (fetch, decode, ...)
 - Modules are ordered along the flow of computation
 - One module's output is the input of the next module
- Turn each module into a pipeline stage
 - Add pipeline registers before every stage except the 1st
 - These store the inputs for that stage
 - Stages execute in parallel working on different instructions
- As we will see later this introduces new problems

PIPELINE STAGES

How many stages should a pipeline have?

- If it has too few stages...
 - We are not exploiting the parallelism present in the program
- If it has too many...
 - There is high overhead and complexity
 - The program may not have enough parallelism to use them well

EXAMPLES

- MIPS processor (1985): first RISC processor, 5 stages
- Sparc, PowerPC processors: 9 pipeline stages
- Intel Pentium IV (late models): _____
- Intel Core i7 processor: _____


MEASURING EFFICIENCY

- **Latency:**
 - How long it takes to execute one instruction from start to finish
 - Will usually not be reduced in a pipeline
- **Throughput:**
 - The number of instructions we can execute per unit of time
 - This is the only meaningful measure for pipelined CPUs

LATENCY & THROUGHPUT


- Sequential Implementation

Minimum times needed
for each stage



- Latency: _____
- Throughput: _____

LATENCY AND THROUGHPUT



Assuming all stages are in use

- Throughput: _____
- Latency: _____

LATENCY AND THROUGHPUT

- Generalizing for pipelined CPUs:
 - Stages require an additional overhead
 - Storing and retrieving special registers
 - Latency for one instruction increases
 - New instruction can start executing once first stage is complete
 - Better throughput overall
 - All stages must run in same time slot
 - Can't move to the next instruction until slowest stage is free

PIPELINED Y86 IMPLEMENTATION

- Unit outline
 - Motivation and basic concepts
 - Initial Implementation
 - Hazards
 - Types of hazards
 - Dealing with hazards by stalling
 - Data hazards: forwarding to avoid stalling
 - Control hazards: branch prediction
 - Indirect jumps
 - Performance analysis

PIPELINED COMPUTATION

- divide execution into modules along flow of computation
 - classic RISC pipeline has five modules
 - modules arranged in order along computation flow
 - all inputs to a module must be computed by an earlier module in flow
- turn modules into pipeline stages
 - add pipeline registers between each stage
 - registers store inputs for that stage
 - each stage executes in parallel working on a different instruction
- observe that
 - a stage has less gate-propagation delay than overall circuit
 - what determines clock rate?

THE RISC PIPELINE

- How many stages?
 - enough to achieve sufficient parallelism
 - must have enough parallelism in the program
 - not too much to add undue overhead or complexity
- Which stages?
 - divide instructions into stages that instruction completes in order
 - then we can execute the stages in parallel on different instructions
- Example: `rmmovq rA, D(rB)`
 - what are the parts?
 - what order do they need to execute?

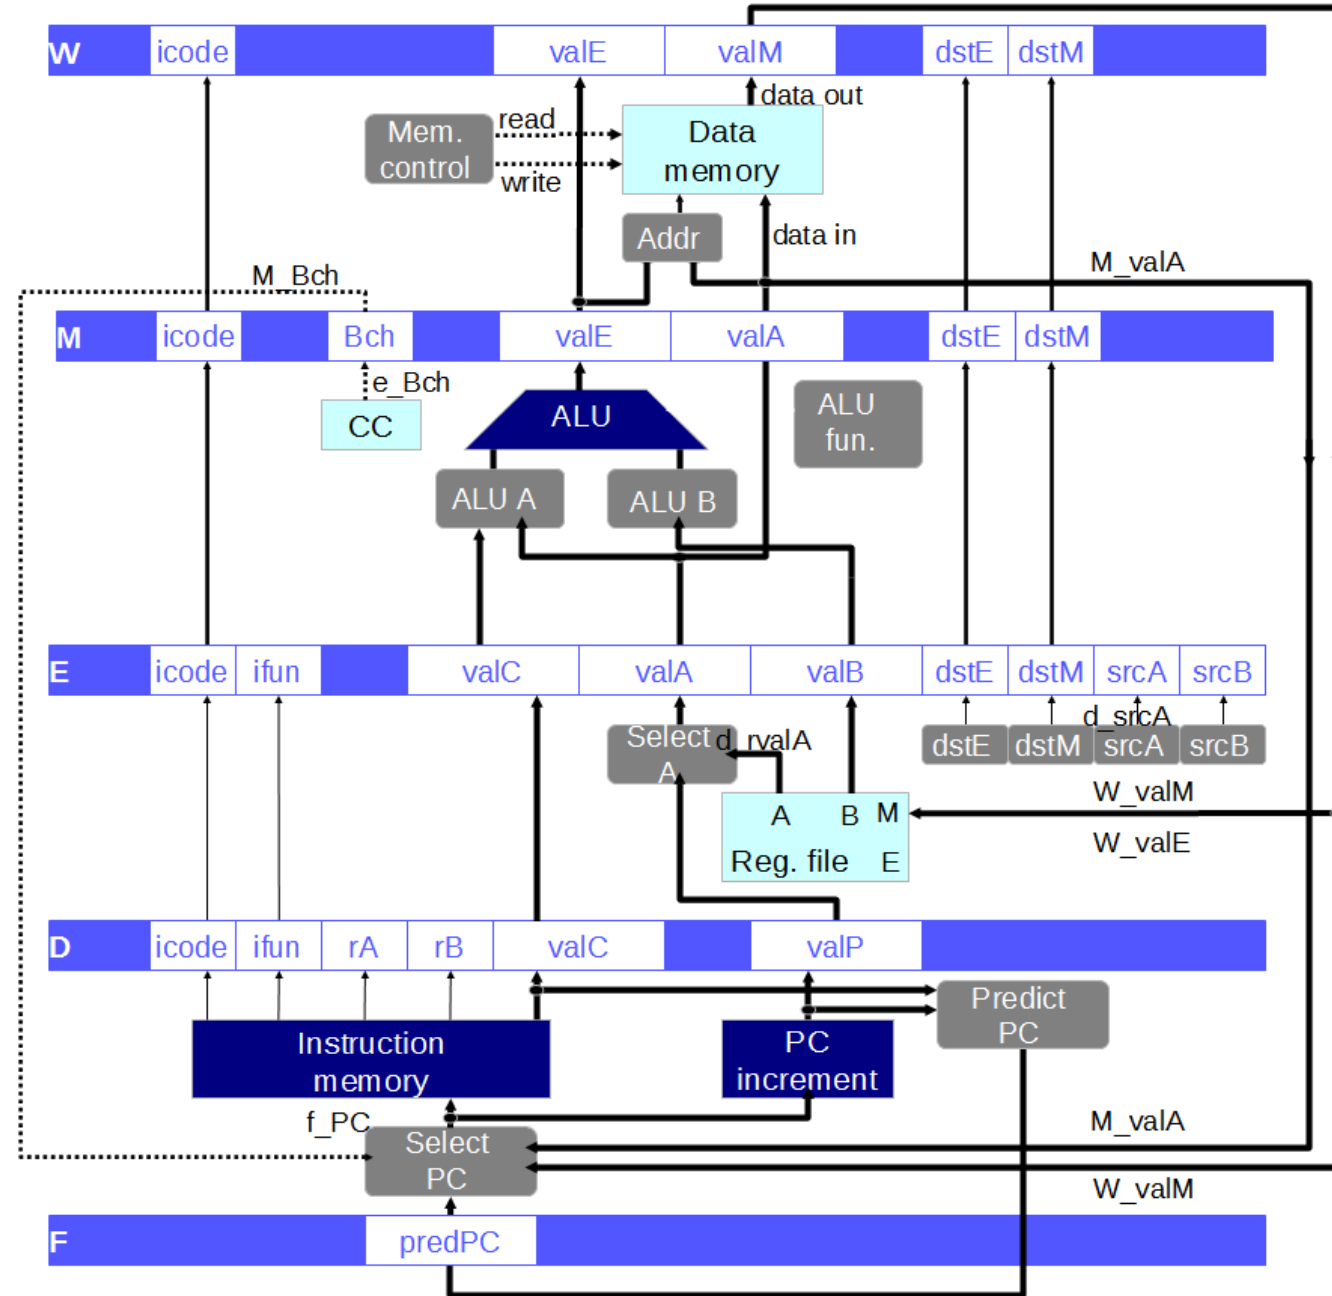
Write back

Memory

Execute

Decode

Fetch



AN EXAMPLE PROBLEM

- What will each pipeline register contain when the last instruction in this sequence is entering the Fetch stage?

```
.pos 0x100
```


```
irmovq $1, %rax
```

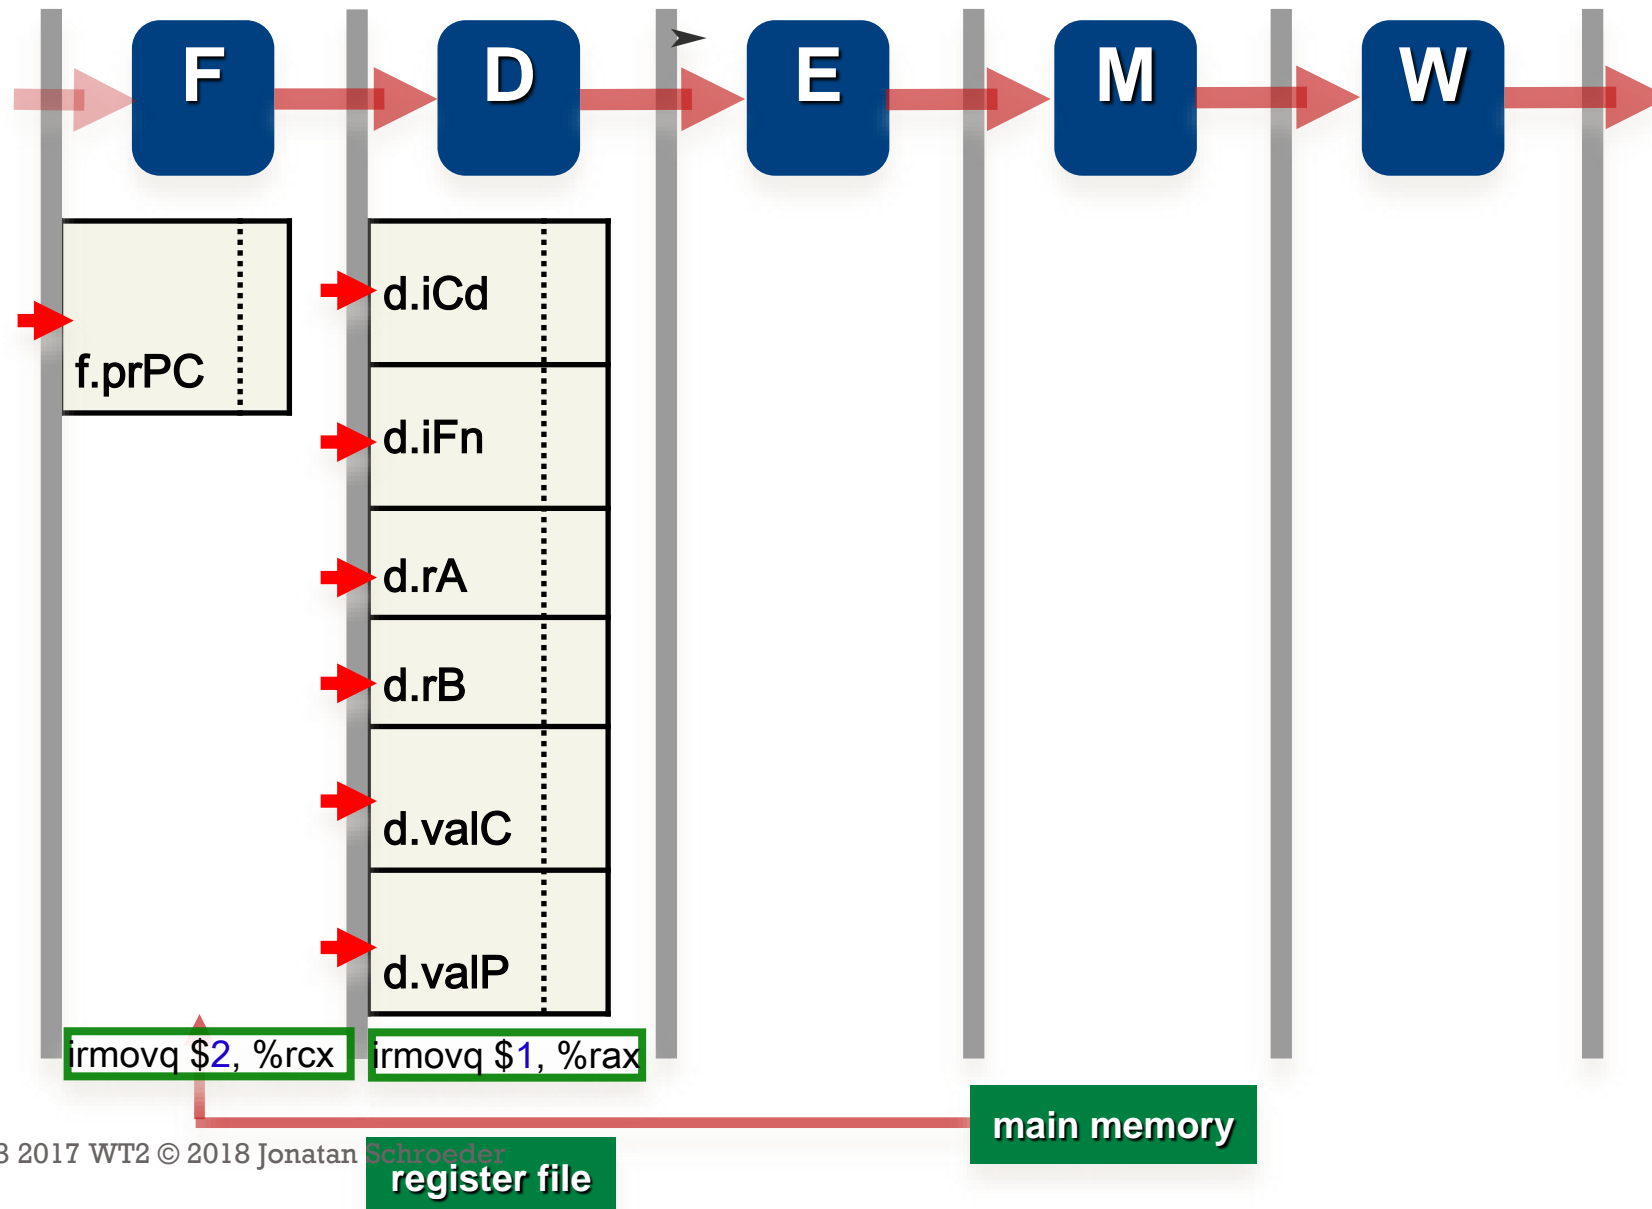
```
irmovq $2, %rcx
```

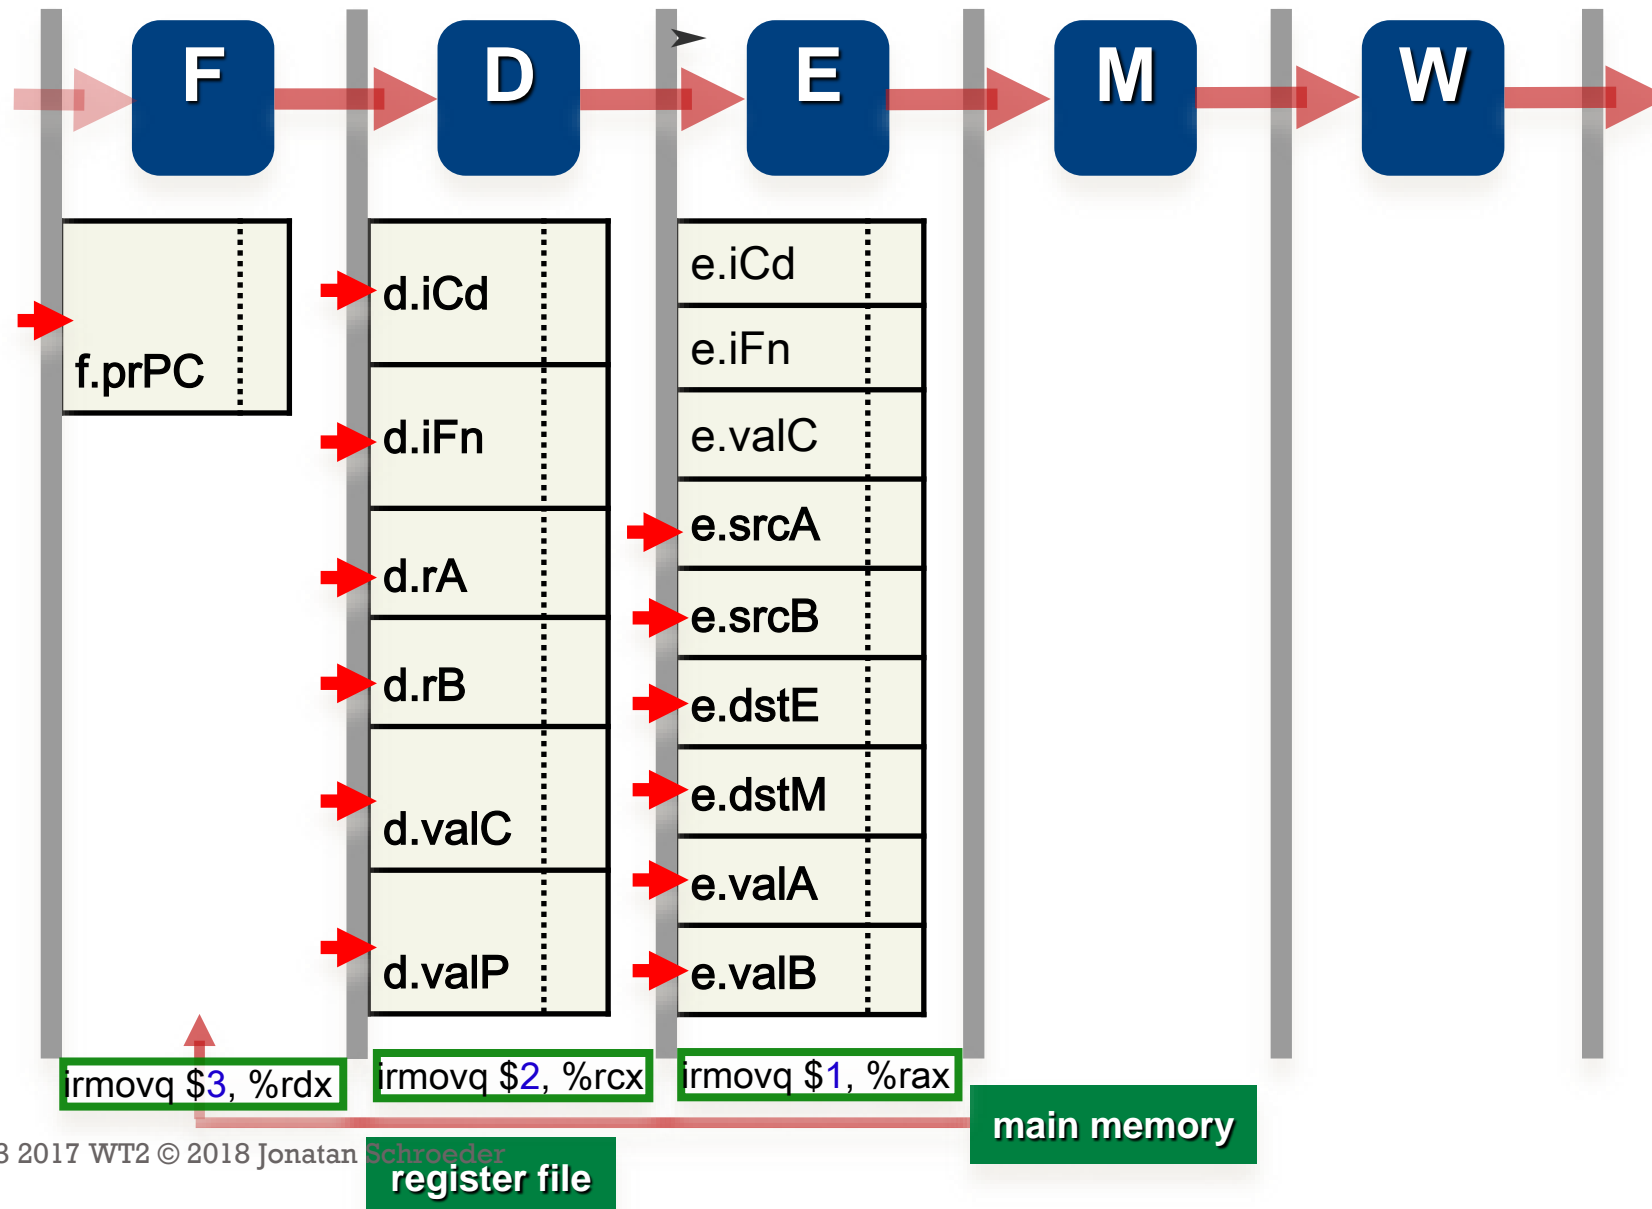
```
irmovq $3, %rdx
```

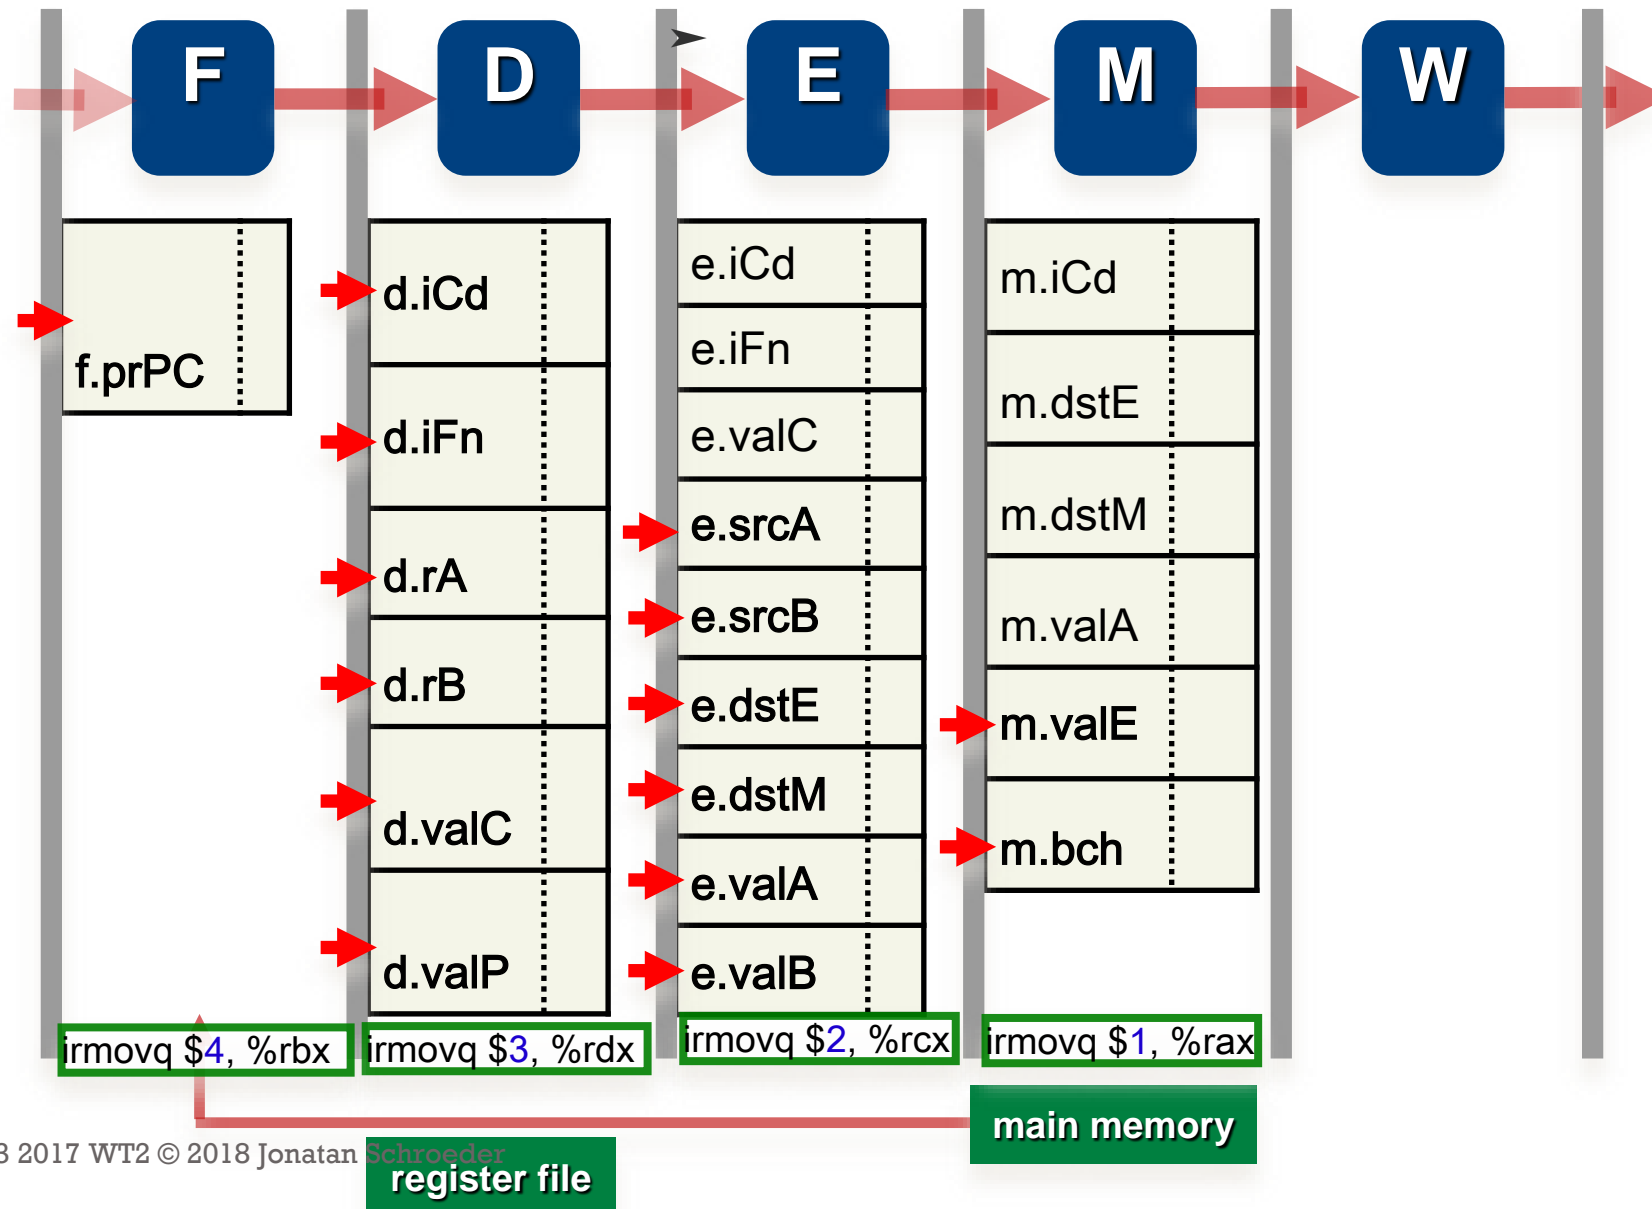
```
irmovq $4, %rbx
```

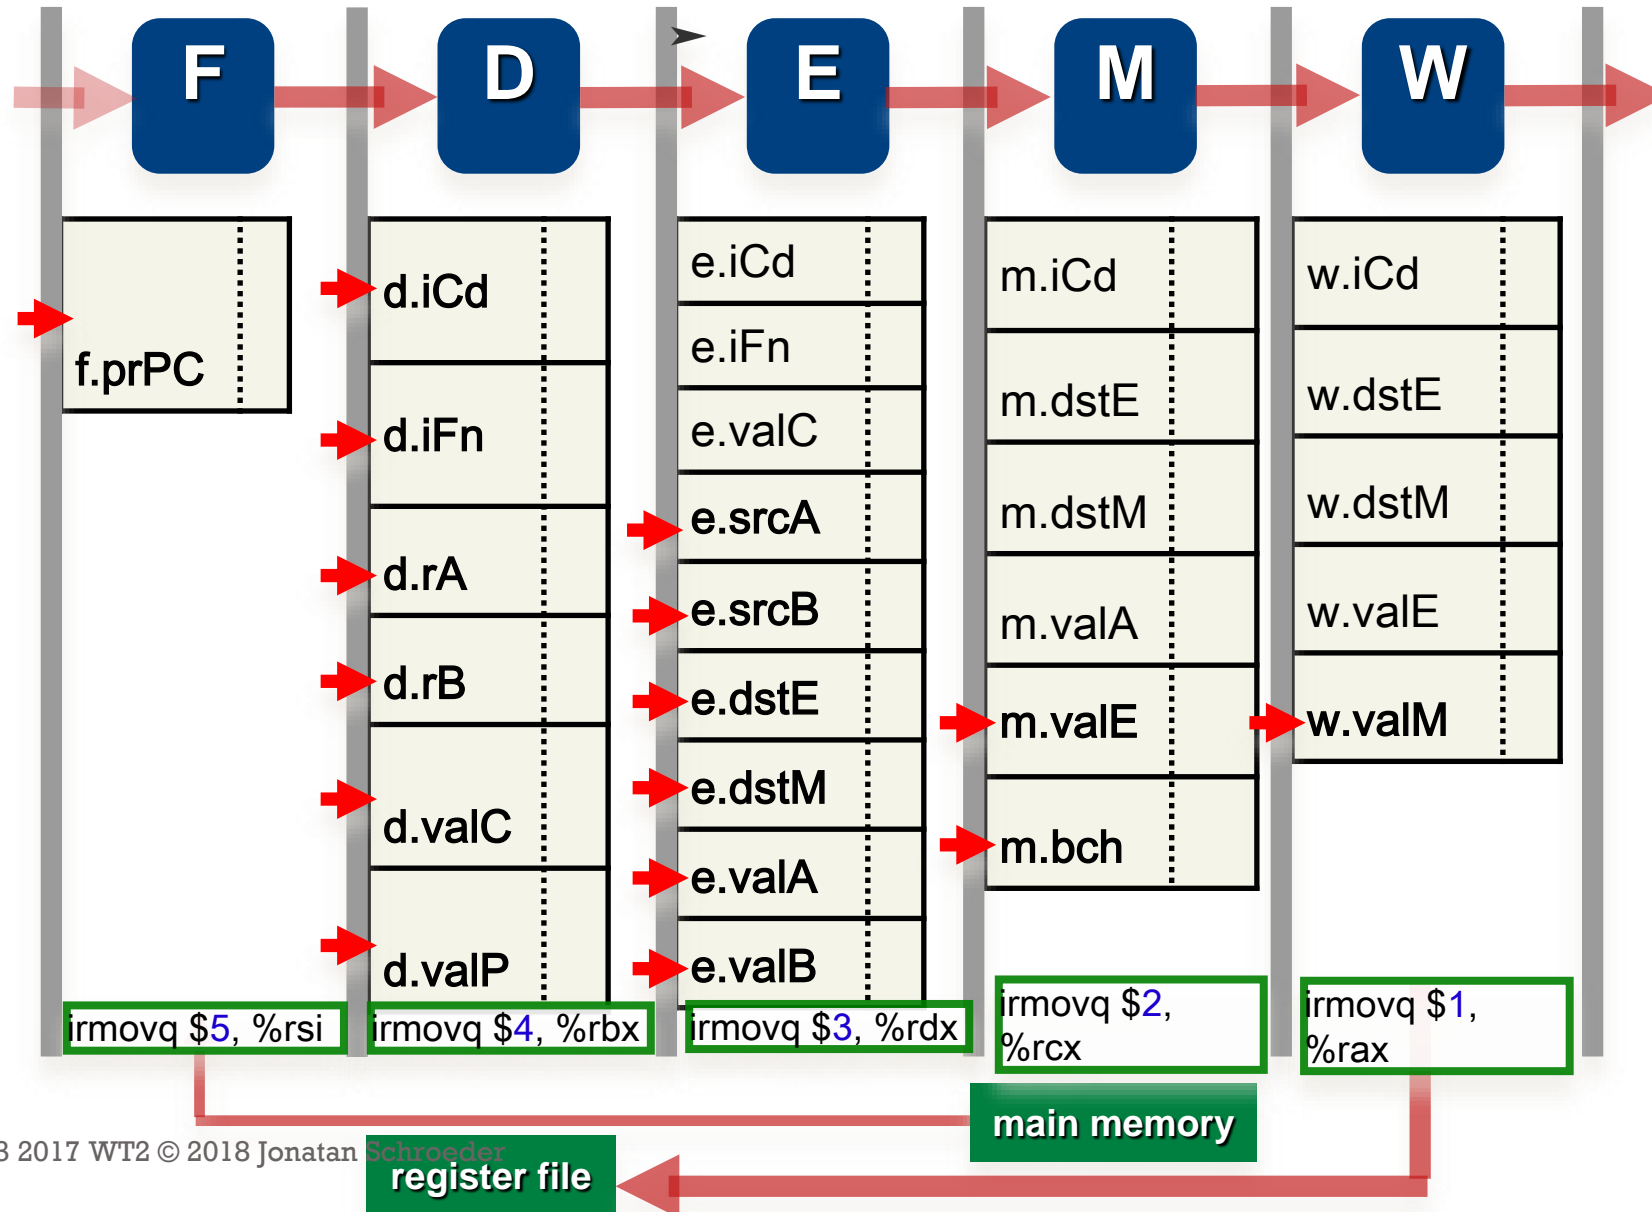
```
irmovq $5, %rsi
```









TERMINOLOGY

- an instruction is *in flight* when
 - it is executing in the pipeline
- an instruction is *retired* when
 - it exits the pipeline; i.e., it completes
- on upcoming slides ...
 - *exploiting* and *expressing* parallelism
 - *instruction-level parallelism*
 - instruction *dependencies*
 - *thread-level parallelism*
 - *sequential consistency*
 - pipeline *hazard, stall and bubble*

EXPRESSING VS EXPLOITING PARALLELISM

- Expressing parallelism: it's a mechanism where the programmer tells the system that two pieces of code can execute in parallel
- Exploiting parallelism: it's the system actually executing two pieces of code in parallel
- How do you express parallelism in C or Java?
- Is this mechanism useful for expressing ILP?

EXPLOITING INSTRUCTION-LEVEL PARALLELISM

The problem with instruction-level parallelism is:

- Programming languages like C, C++ and Java are based on the sequential consistency model:
 - The effect of executing the program must be the same as if instructions were executed one by one in the order they are written
- Programmers write code without thinking about parallelism
 - Example:
a = b + c; d = a + 1; e = f + g;
Can be rewritten as:
a = b + c; e = f + g; d = a + 1;
- Compilers must find this on their own

PIPELINED Y86 IMPLEMENTATION

- Unit outline
 - Motivation and basic concepts
 - Initial Implementation
 - Hazards
 - Types of hazards
 - Dealing with hazards by stalling
 - Data hazards: forwarding to avoid stalling
 - Control hazards: branch prediction
 - Indirect jumps
 - Performance analysis

PIPELINE CONSIDERATIONS

- Consider this code:

```
irmovq $5, %rax
```

```
addq    %rax, %rdx
```

- At what stage is `irmovq` in when `addq` needs `%rax`?
 - At what stage is the value of `%rax` needed in `addq`?
 - At what stage is the value of `%rax` updated in `irmovq`?

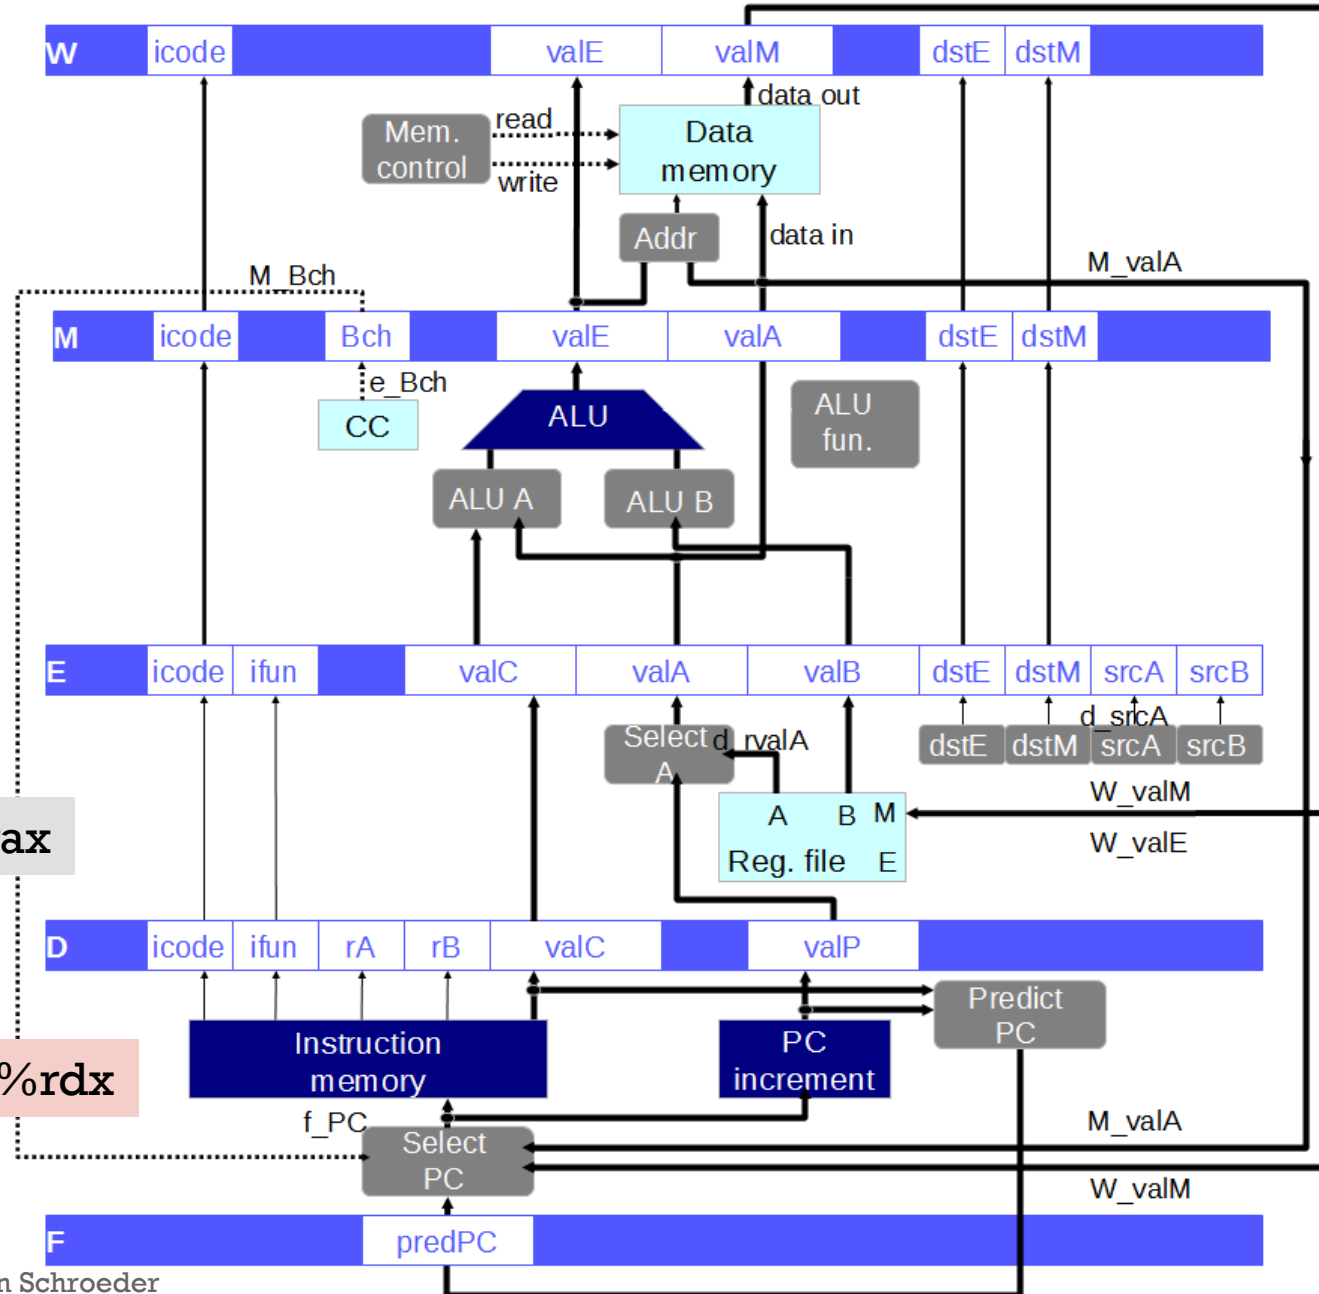
Write back

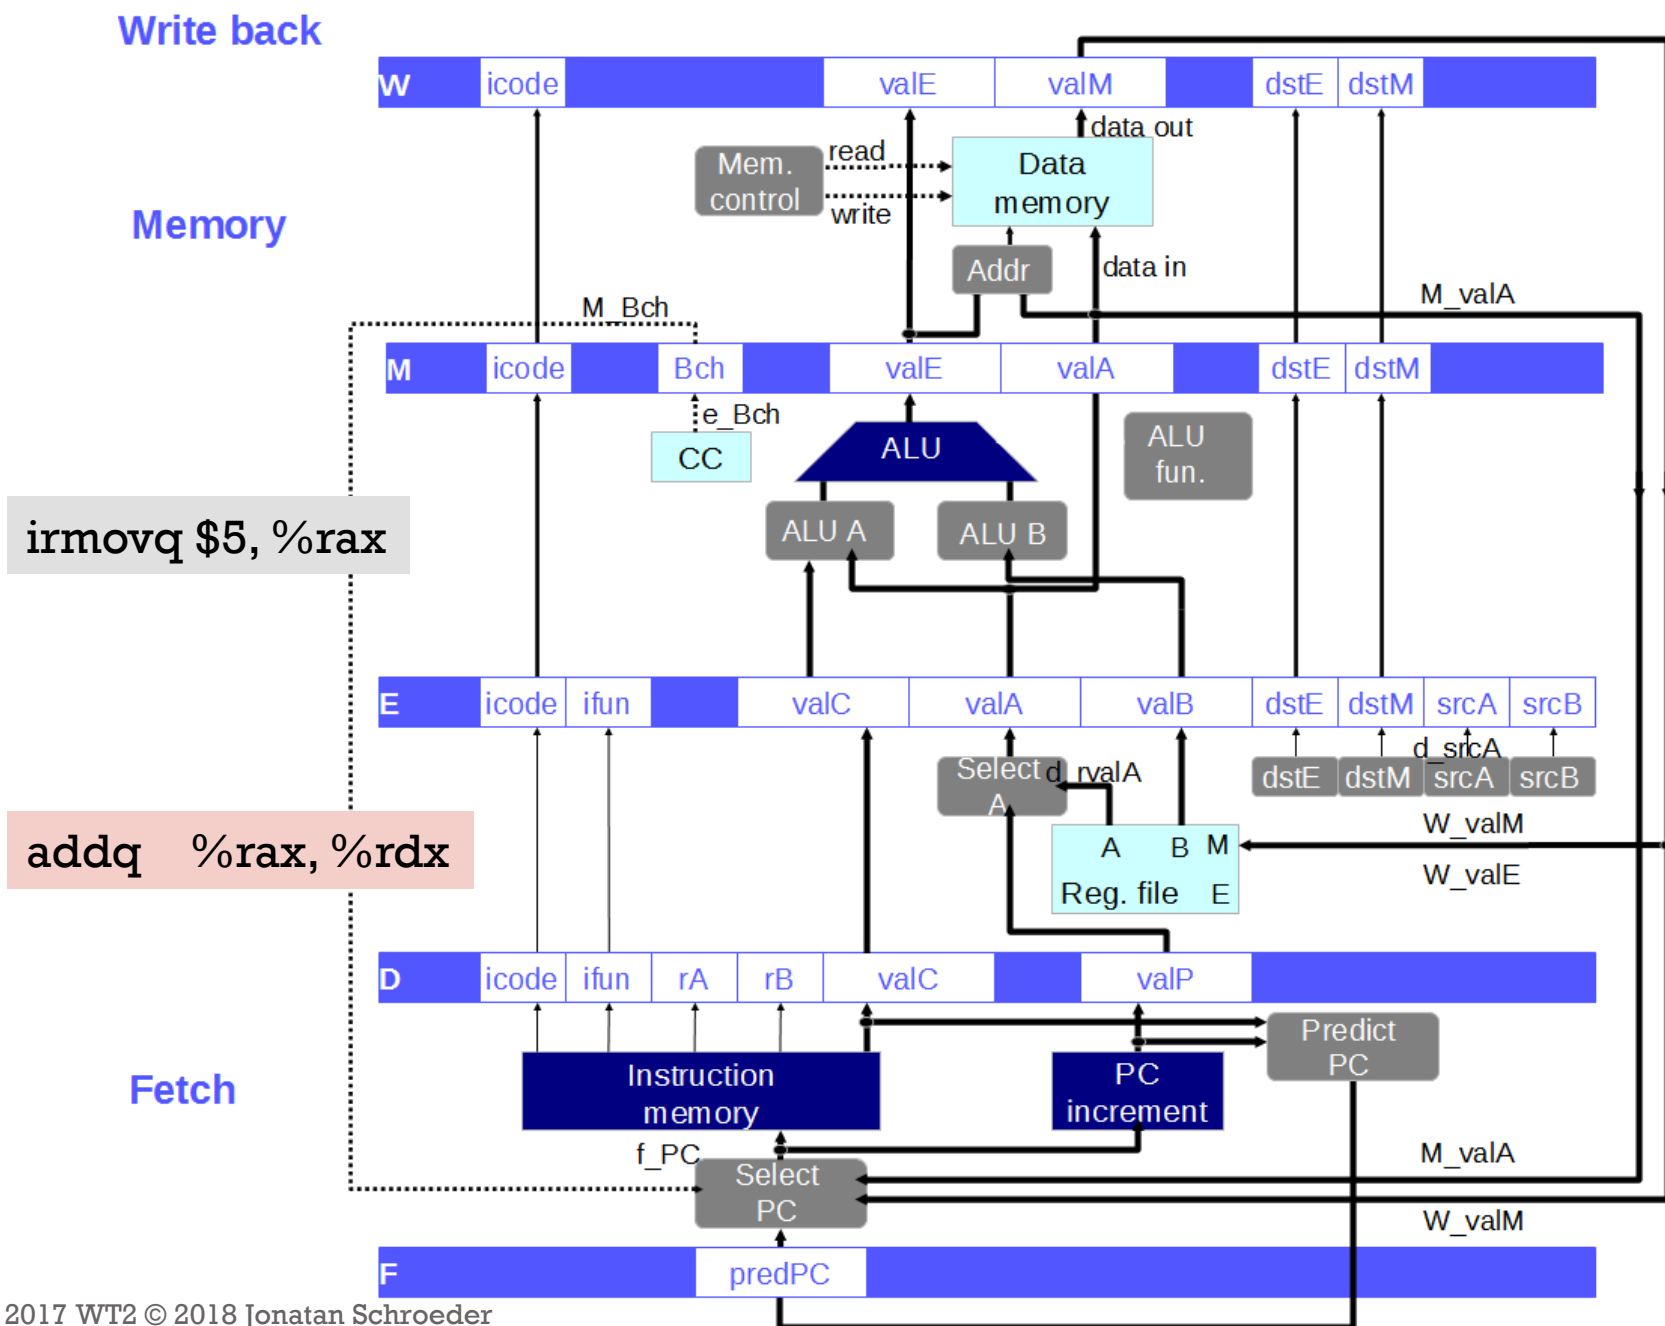
Memory

Execute

`irmovq $5, %rax`

`addq %rax, %rdx`





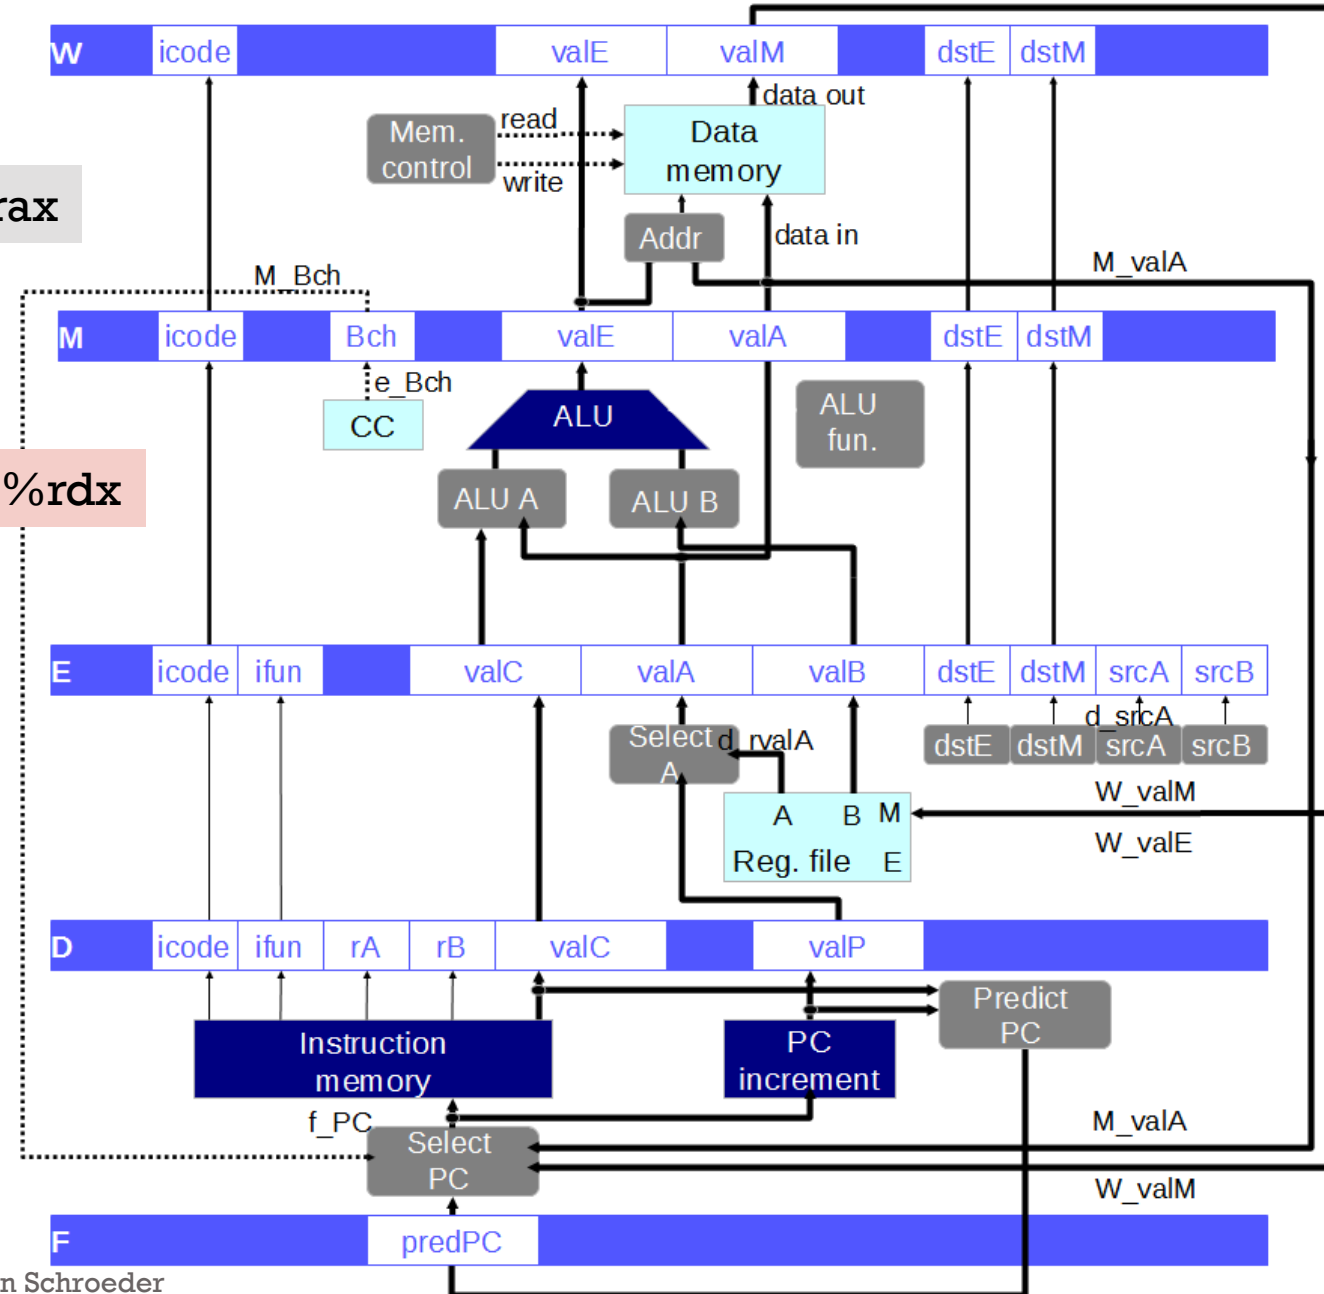
Write back

`irmovq $5, %rax`

`addq %rax, %rdx`

Decode

Fetch



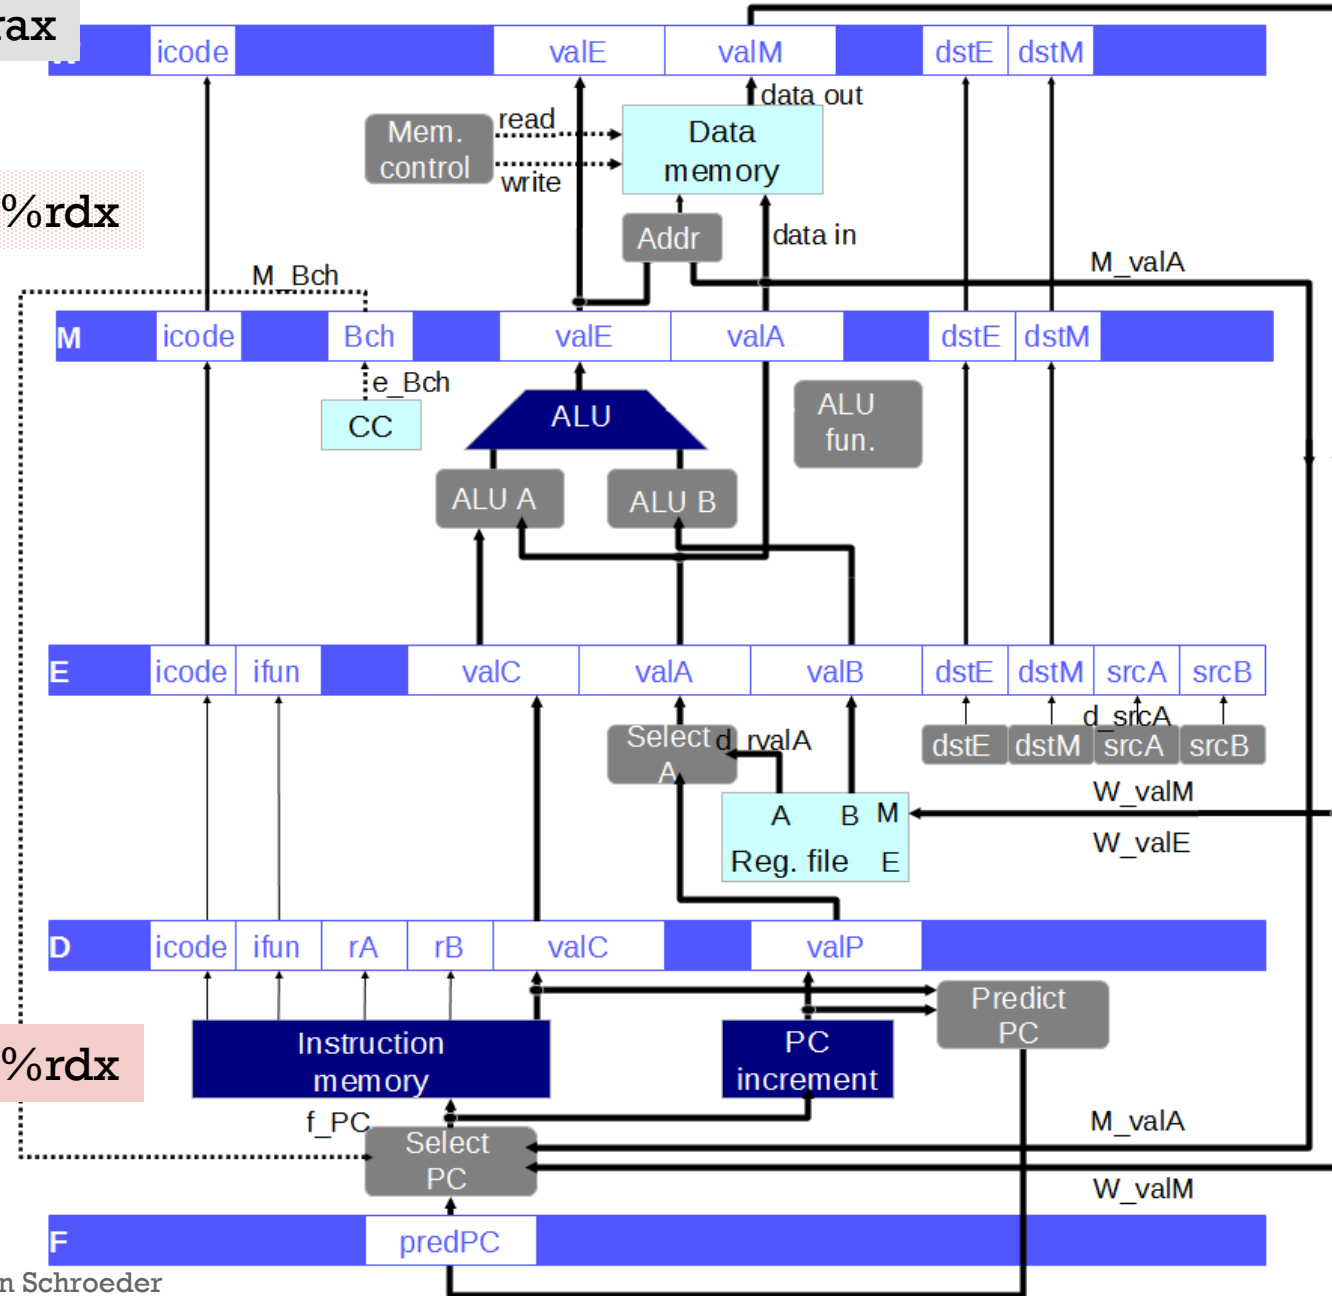
```
irmovq $5, %rax
```

```
addq %rax, %rdx
```

```
addq %rax, %rdx
```

Execute

Decode



INSTRUCTION-LEVEL PARALLELISM

- Pipeline requires some parallelism
 - multiple in-flight instructions run partly at the same time
 - may even run out-of-order
- However, the program **must** have the same output as if instructions were executed sequentially.

DEPENDENCIES CONSTRAIN PARALLELISM

- Execution of 2 or more instructions has to proceed in strict time order
 - Must be able to produce the same output as if one instruction were completely executed before the next instruction is started
- If no dependency, execution order doesn't matter
- Compilers try to expose as much instruction-level parallelism as they can
 - By separating dependent instructions
 - By grouping them with others on which they don't depend

DEPENDENCIES

- Example: how can we rewrite the following code to expose more parallelism?

```
addq %rax, %rbx  
iaddq $5, %rbx  
addq %rdx, %rcx  
iaddq $9, %rcx  
addq %rsi, %rdi  
iaddq $3, %rdi
```

COMPILERS & INSTRUCTION PARALLELISM


- Compilers can reorder instructions to expose as much instruction-level parallelism as possible.
- However they cannot know every detail of the processor's pipeline (e.g. later Pentium IV's had more stages than earlier ones).
 - How many instructions should be kept between dependencies?
- So the **pipeline** must handle all dependencies correctly

DEPENDENCY TYPES

- Data dependencies
 - Involve dependencies between registers
 - Example: one instruction needs a register that is written by a previous instruction
- Control dependencies
 - Involve the flow of control (changes in PC)
 - Example: conditional jumps

CLASSIFYING DATA DEPENDENCIES

- if A and B are instructions, then $A < B$ means instruction B depends on instruction A
 - Causal: $A < B$ if B reads a value written by A
 - Output: $A < B$ if B writes to a location written by A
 - Alias (anti): $A < B$ if B writes to a location read by A



Read “ $<$ ” as must
happen before

TYPES OF DATA DEPENDENCY: EXERCISE

a) no dependency

b) casual dependency

c) anti-dependency

d) output dependency

1.

a = 1;

b = 2;

3.

a = 1;

b = a;

2.

a = 1;

a = 2;

4.

a = b;

b = 2;

TYPES OF DATA DEPENDENCY: EXERCISE

a) no dependency

c) anti-dependency

b) casual dependency

d) output dependency

1.

```
irmovq $1, %rax  
rrmovq %rcx, %rax
```

3.

```
rrmovq %rax, %rcx  
irmovq $1, %rax
```

2.

```
irmovq $1, %rax  
rrmovq %rax, %rcx
```

4.

```
irmovq $1, %rax  
irmovq $2, %rcx
```

CONTROL DEPENDENCIES

- Control dependencies determine what code is executed next
- Examples:
 - Whether a branch is taken or not taken (e.g., conditional jumps)
 - When the next instruction is obtained from a register or memory (e.g., return)
 - When an instruction writes to instruction memory (self-modifying code)
 - Will not be explored

WHEN DEPENDENCIES BECOME HAZARDS

- Dependencies are not always a problem
 - If there are enough instructions in between dependent instructions, there is no hazard
- Hazard happens when “normal” pipeline execution violates the dependency
- CPU must change its behaviour
 - By stalling instructions
 - By forwarding values from previous instructions
 - By speculating what instructions must run

PIPELINED Y86 IMPLEMENTATION

- Unit outline
- Motivation and basic concepts
- Initial implementation
- Hazards
 - Types of hazards
 - Dealing with hazards by stalling
 - Data hazards: using data forwarding to avoid stalling
 - Control hazards: branch prediction
 - Indirect jumps
- Performance analysis.

LEARNING GOALS

- Explain how a pipelined processor uses stalling, data forward, and branch prediction to reduce or eliminate hazards.
- Define what is meant by a pipeline bubble and explain why/how a bubble is generated.
- For a sequence of machine language instructions, describe at an arbitrary execution point the state of the pipeline:
 - When there are no hazards
 - When hazards are handled using only stalling
 - When hazards are handled using stalling and/or data forwarding
 - When hazards are handled using stalling, data forwarding, and/or branch prediction

WHAT IS THE PROBLEM HERE?


`irmovq $1, %rax`

`addq %rax, %rbx`

`addq %rax, %rcx`

`addq %rax, %rdx`

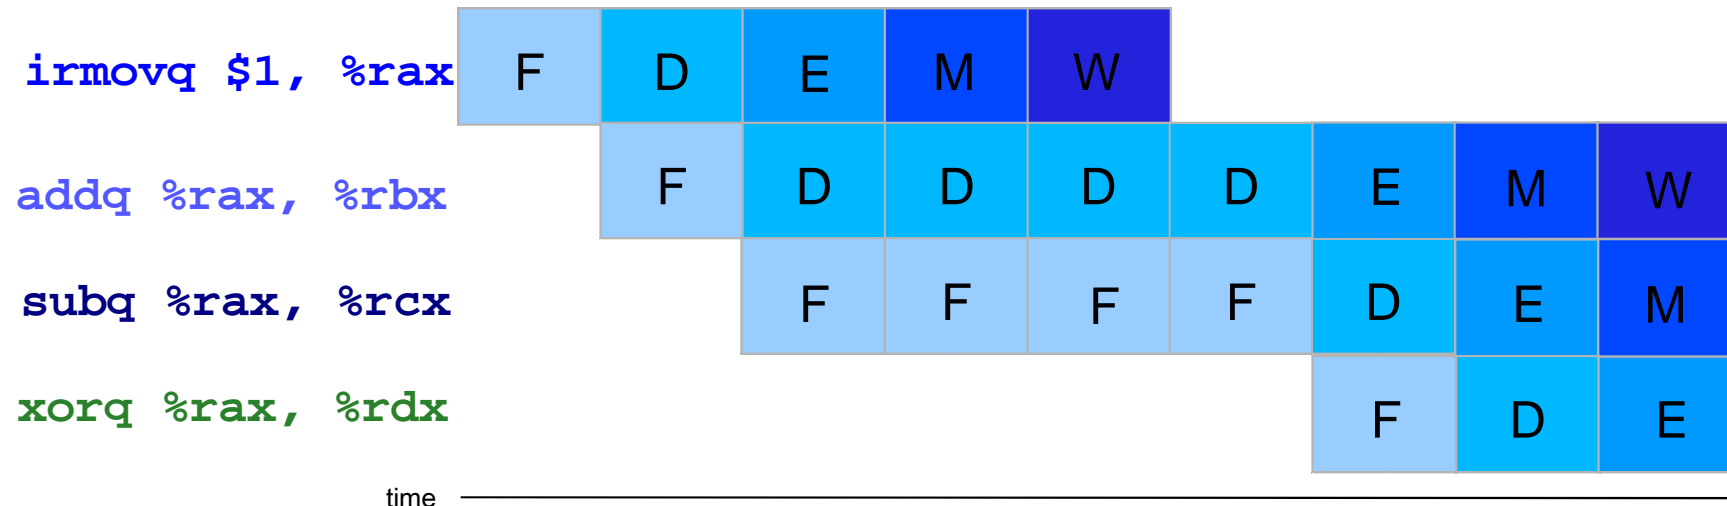
`addq %rax, %rsi`



time

STALLING

- pipeline stall: hold an instruction in a pipeline stage for an extra cycle



BUBBLE

- The term **pipeline bubble** denotes a pipeline stage that is forced to do nothing to avoid a hazard, because of a stall in a previous stage
 - For example, if decode stalls, in the next instruction execute will have a bubble
- The bubble converts the instruction into a NOP

EXAMPLE

Fetch	Decode	Execute	Memory	Write-back
irmovq	?	?	?	?
addq	irmovq	?	?	?

EXAMPLE

Fetch	Decode	Execute	Memory	Write-back
irmovq	?	?	?	?
addq	irmovq	?	?	?
subq	addq	irmovq	?	?

EXAMPLE

Fetch	Decode	Execute	Memory	Write-back
irmovq	?	?	?	?
addq	irmovq	?	?	?
subq	addq	irmovq	?	?
subq	addq	bubble	irmovq	?

EXAMPLE

Fetch	Decode	Execute	Memory	Write-back
irmovq	?	?	?	?
addq	irmovq	?	?	?
subq	addq	irmovq	?	?
subq	addq	bubble	irmovq	?
subq	addq	bubble	bubble	irmovq

EXAMPLE

Fetch	Decode	Execute	Memory	Write-back
irmovq	?	?	?	?
addq	irmovq	?	?	?
subq	addq	irmovq	?	?
subq	addq	bubble	irmovq	?
subq	addq	bubble	bubble	irmovq
subq	addq	bubble	bubble	bubble

EXAMPLE

Fetch	Decode	Execute	Memory	Write-back
irmovq	?	?	?	?
addq	irmovq	?	?	?
subq	addq	irmovq	?	?
subq	addq	bubble	irmovq	?
subq	addq	bubble	bubble	irmovq
subq	addq	bubble	bubble	bubble
xorq	subq	addq	bubble	bubble

EXAMPLE

Fetch	Decode	Execute	Memory	Write-back
irmovq	?	?	?	?
addq	irmovq	?	?	?
subq	addq	irmovq	?	?
subq	addq	bubble	irmovq	?
subq	addq	bubble	bubble	irmovq
subq	addq	bubble	bubble	bubble
xorq	subq	addq	bubble	bubble
	xorq	subq	addq	bubble

EXAMPLE

Fetch	Decode	Execute	Memory	Write-back
irmovq	?	?	?	?
addq	irmovq	?	?	?
subq	addq	irmovq	?	?
subq	addq	bubble	irmovq	?
subq	addq	bubble	bubble	irmovq
subq	addq	bubble	bubble	bubble
xorq	subq	addq	bubble	bubble
	xorq	subq	addq	bubble
		xorq	subq	addq

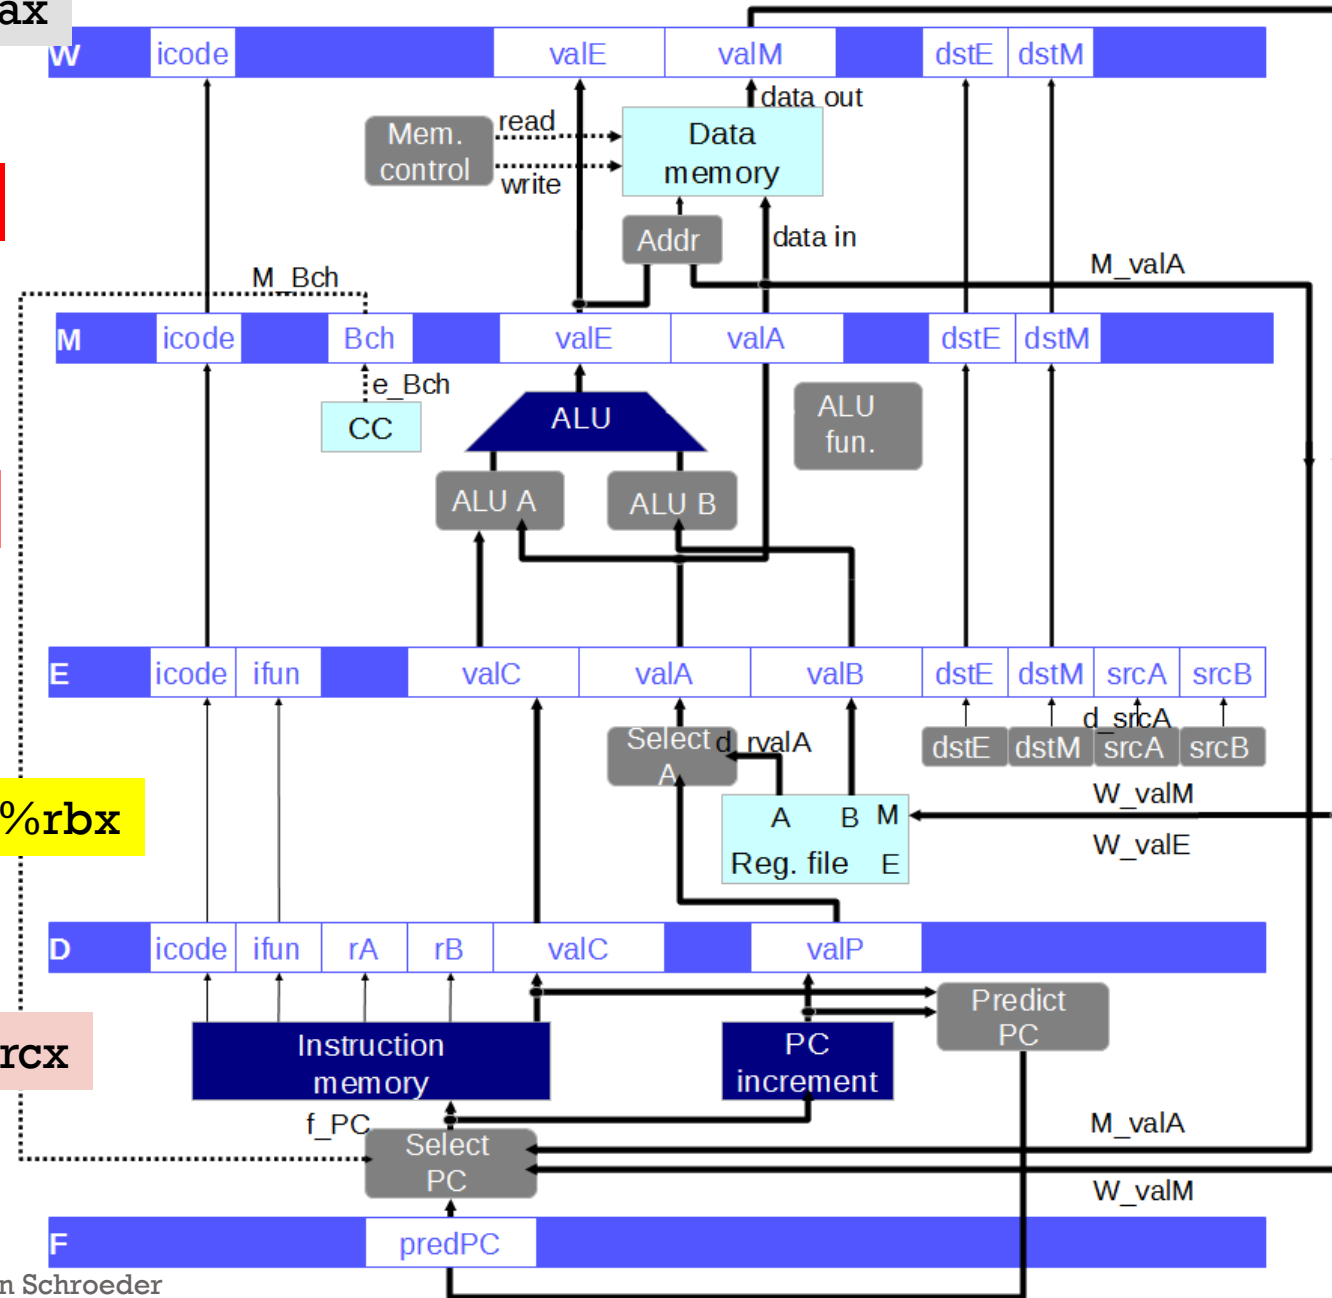
irmovq \$5, %rax

Bubble/ NOP

Bubble/ NOP


addq %rax, %rbx

subq %rax, %rcx



WHAT ABOUT ANTI AND OUTPUT DEPENDENCIES

- If instructions are executed in-order, they are not a problem
- Output dependency: first instruction's output can be ignored




TYPES OF HAZARDS

For our pipelined Y86 implementation:

Anti-dependency:

`addq %rax, %rdx`


`irmovq $1, %rax`



Output dependency:

`addq %rax, %rdx`

`irmovq $1, %rdx`



As long as instructions execute in order, no problem

STALLING WORKS, BUT...

- Stalls should be avoided
 - every stall creates a pipeline bubble
 - every bubble results in a cycle when processor can't retire an instruction
 - retiring a bubble is of no value to program execution
- Bubbles decrease overall throughput
- Ideas needed to deal with
 - data dependencies
 - control dependencies

PIPE: RESOLVING HAZARDS BY STALLING

SUMMARY

- data hazards
 - reading registers in decode that are written by instructions currently in execute, memory or write-back stages
 - stall instruction in decode until writer is retired
 - how many stall cycles? _____

PIPE: RESOLVING HAZARDS BY STALLING

SUMMARY (CONT.)

- Conditional jump
 - At what stage do we know which PC?
 - At what stage is the PC needed?
 - how many stall cycles? _____
- Return
 - The next PC is available at the end of?
 - how many stall cycles? _____

PIPELINE CONTROL UNIT

- The pipeline-control module
 - is a hardware component (circuit) separate from the 5 stages
 - examines values across every stage
 - decides whether stage should stall or bubble

DEALING WITH HAZARDS

- Each stage register has a control input that determines what happens when the clock ticks:
 - Normal: register's new value is the input value.
 - Stall: register's new value is the same as its current value.
 - Bubble: register's new value is the same as for NOP.

PIPELINE CONTROL UNIT

